



CernVM Virtual Workshop 2021  
1-2 February 2021  
NIKHEF (remote)

# ESSI

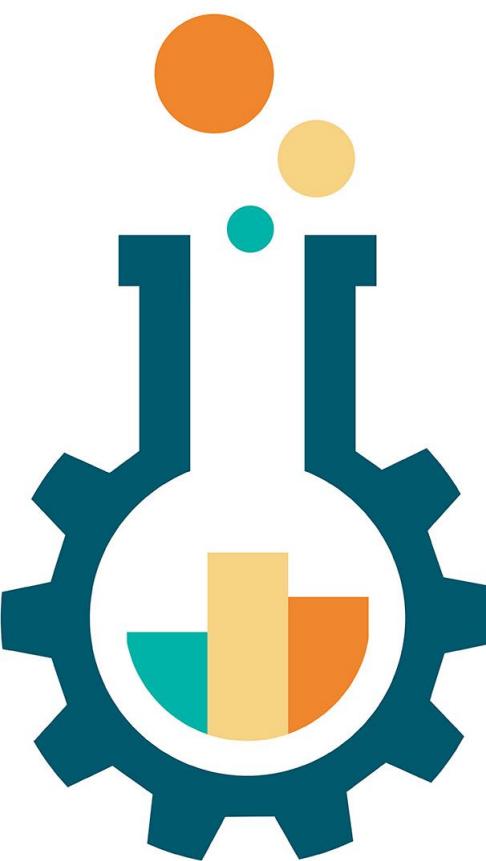
EUROPEAN ENVIRONMENT FOR  
SCIENTIFIC SOFTWARE INSTALLATIONS

*European Environment for Scientific Software Installations*

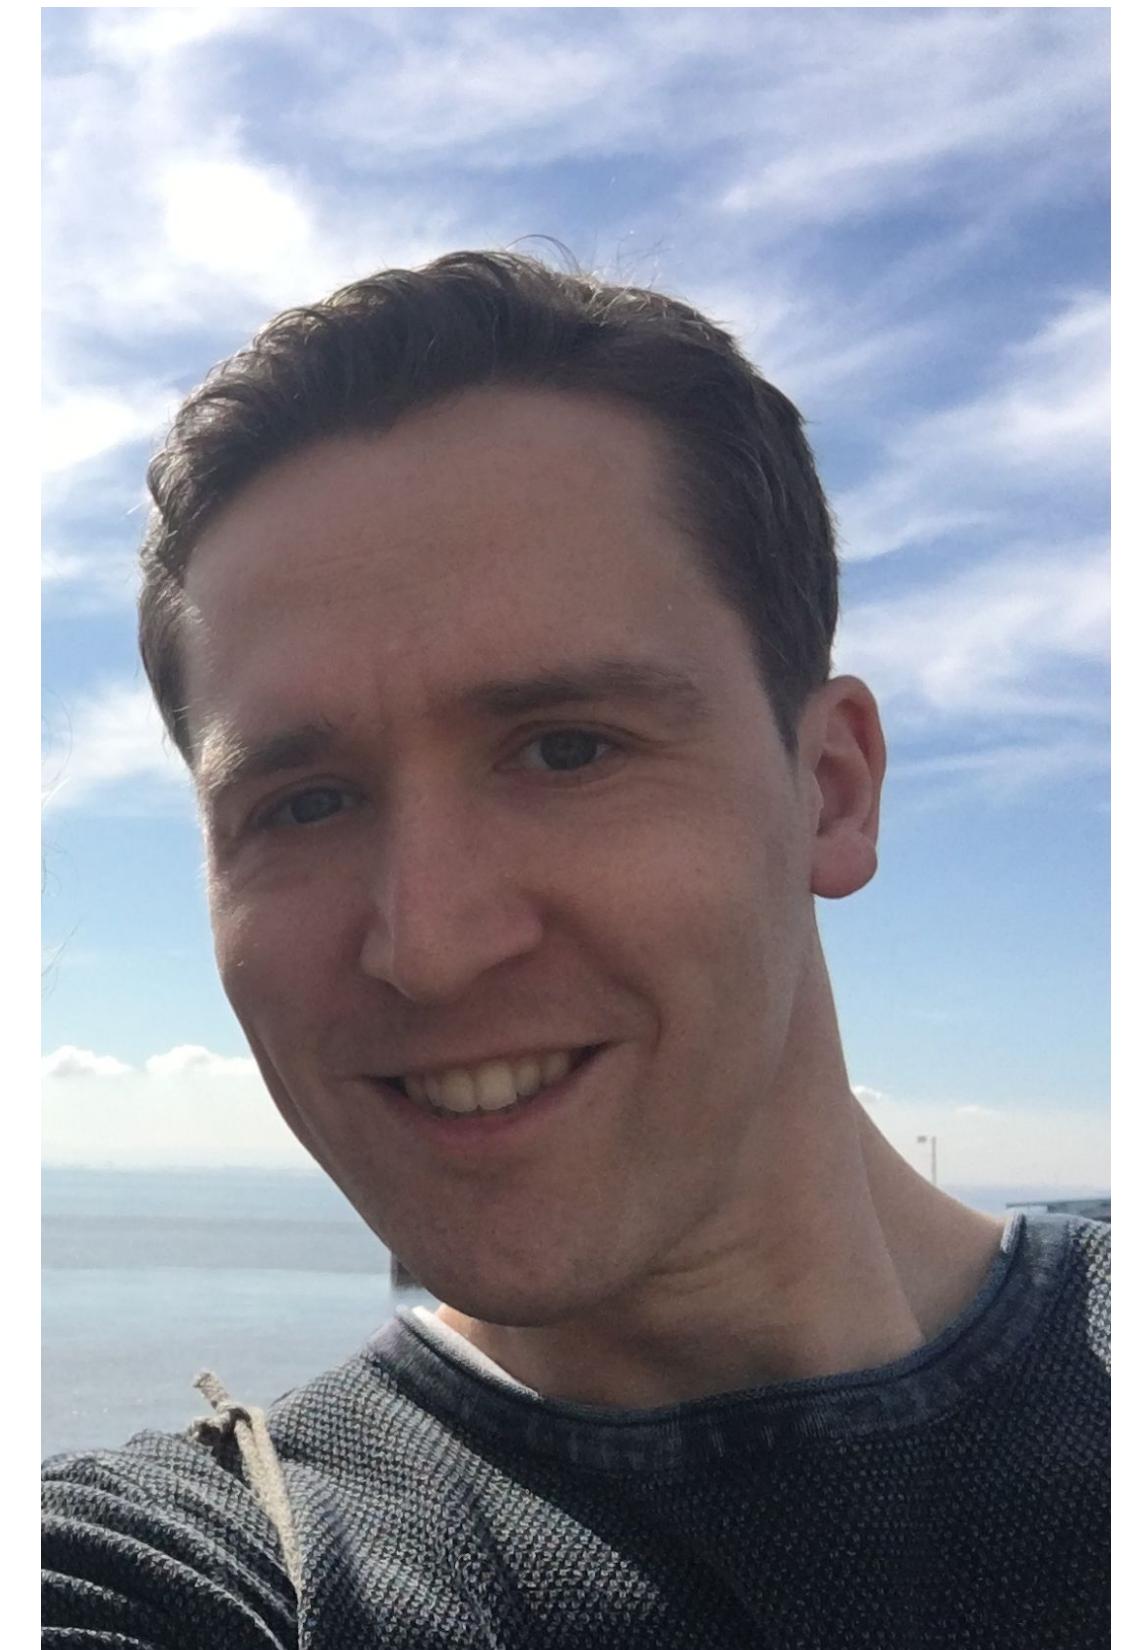
**Bob Dröge, University of Groningen, The Netherlands**

# About me

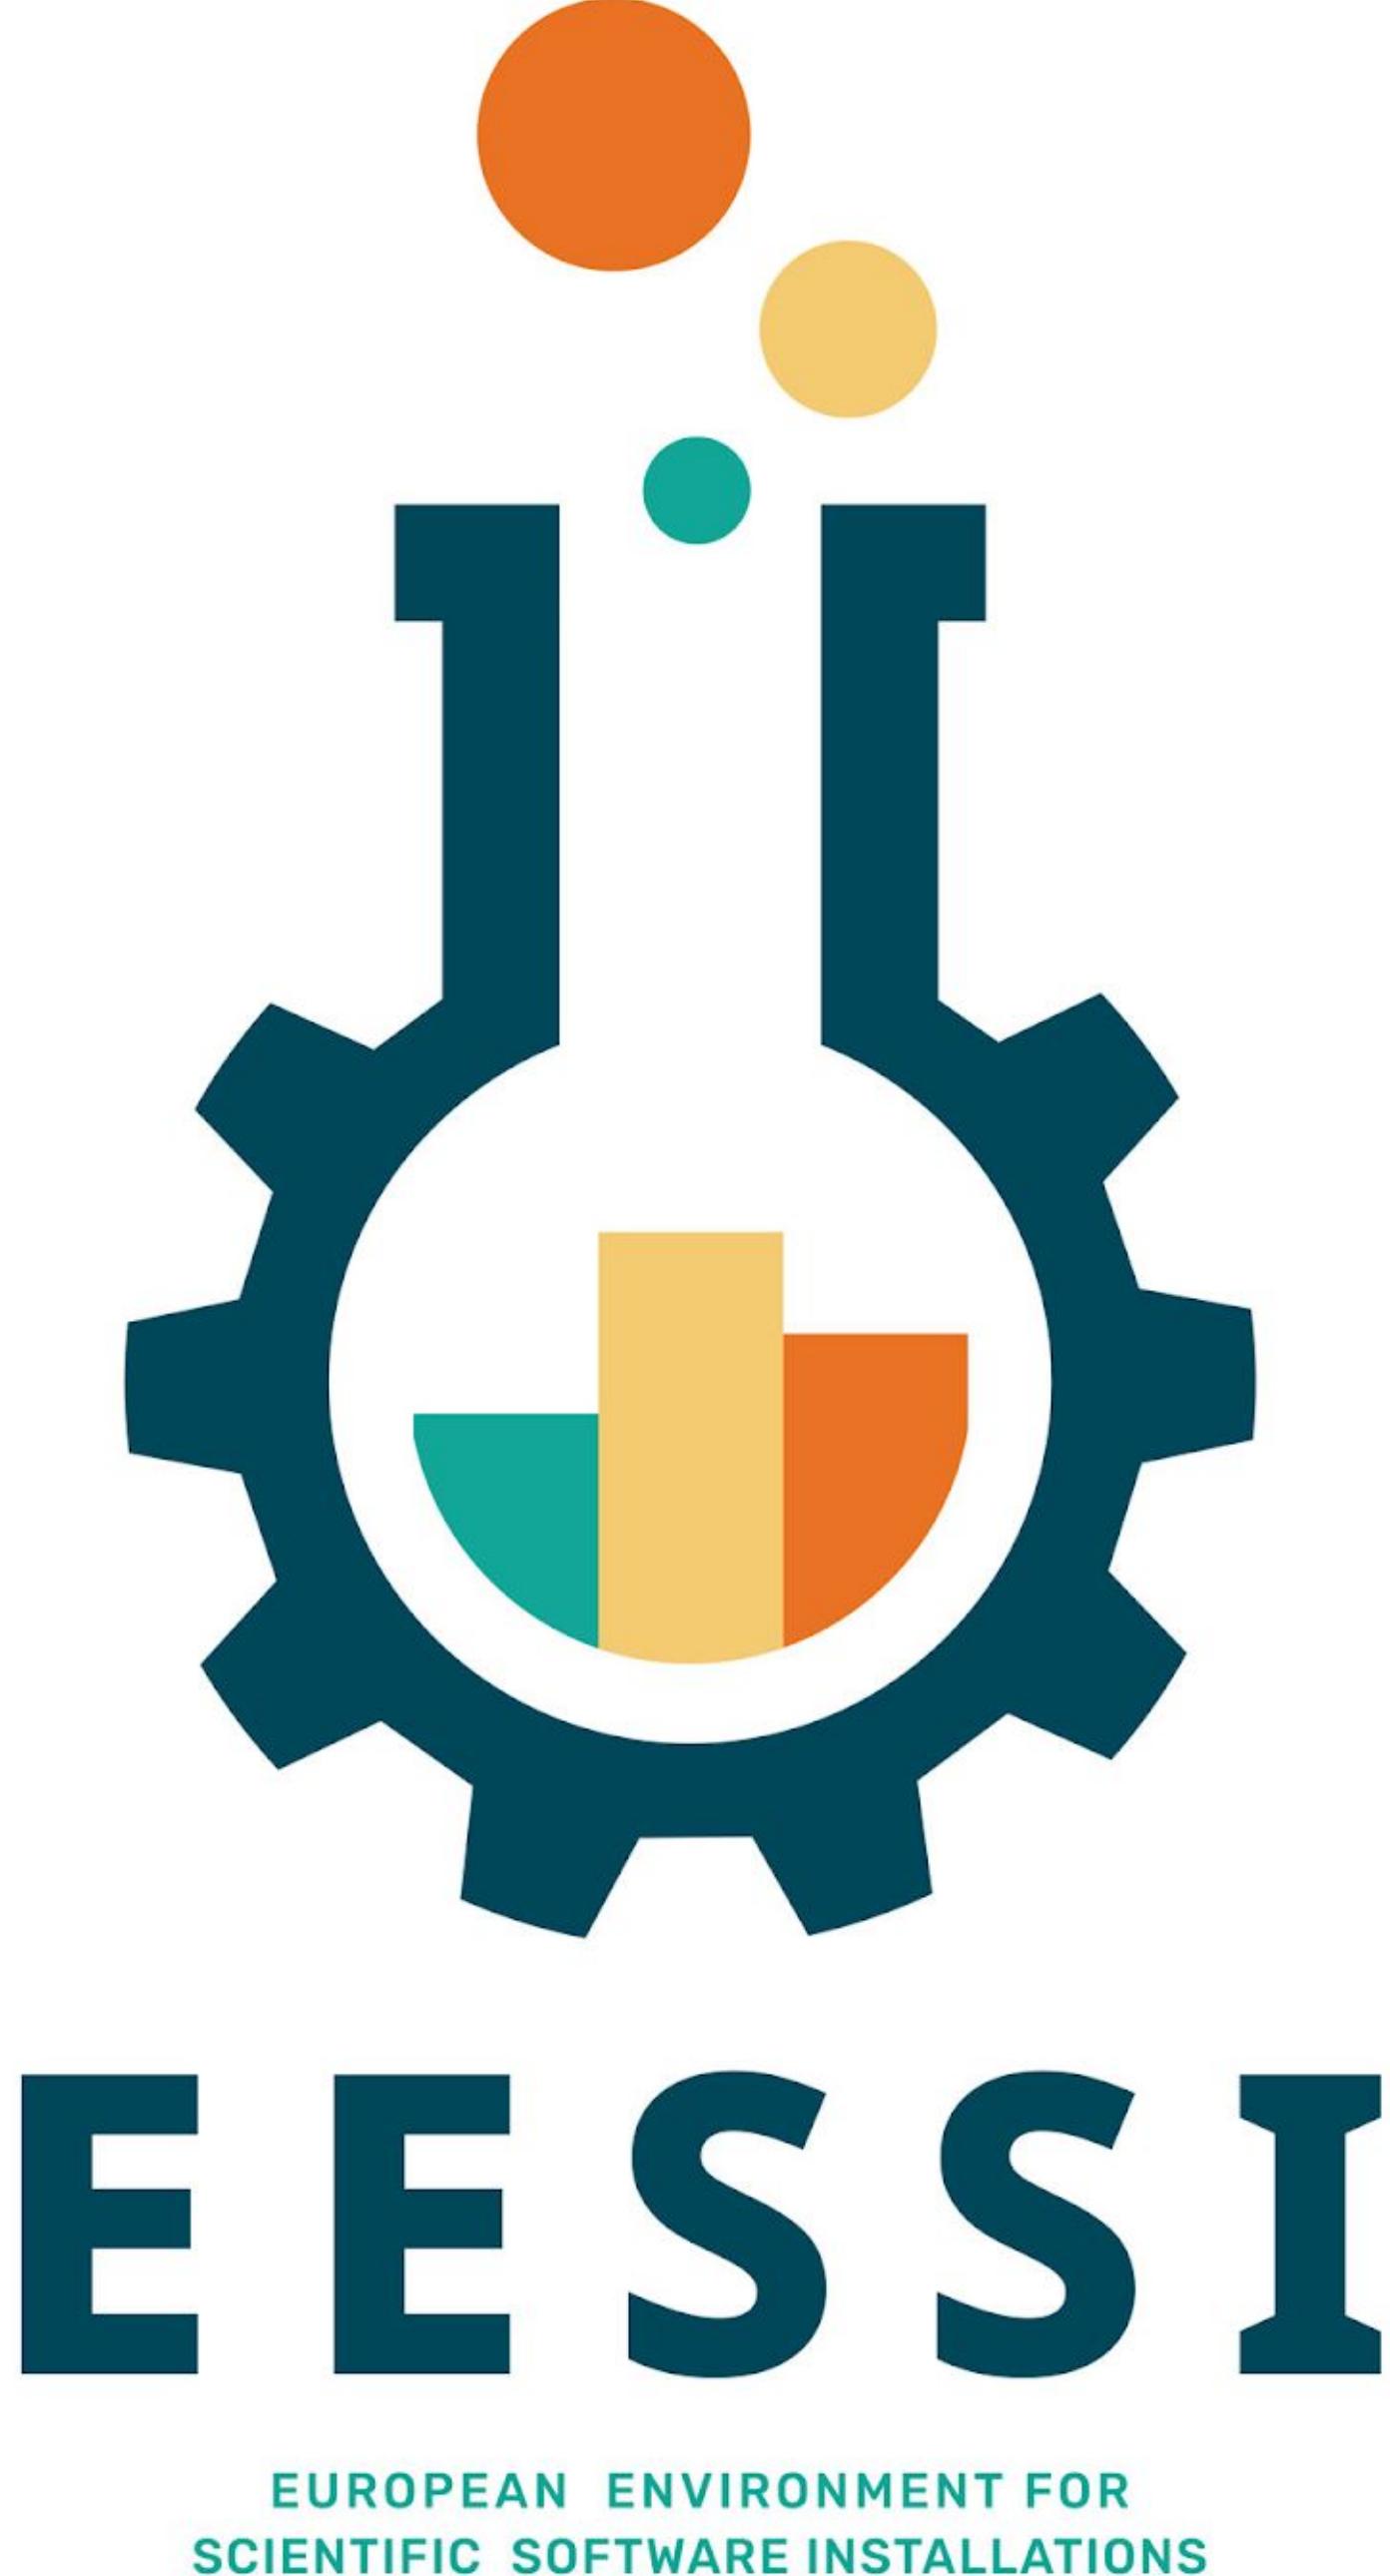
EESSI



- Team High Performance Computing
  - Center for Information Technology
  - University of Groningen
- 
- HPC user support and training,  
installing software
  - HPC system administration
  - Currently involved in two large projects:  
Euclid (ESA project/mission) and EESSI



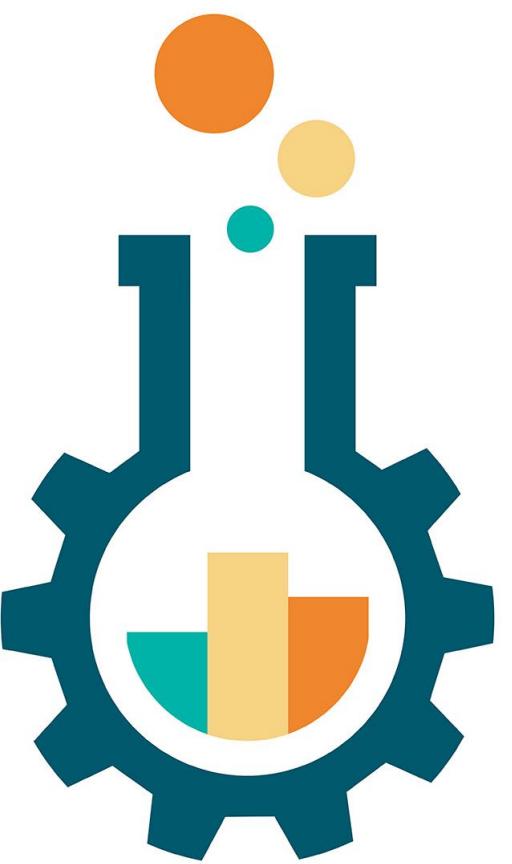
university of  
groningen



- **What** is the project about?
- **Who** is involved in EESSI?
- **Why** did we start it?
- **How** are we tackling the problem?
- Which **FOSS projects** do we use?
- What is the **current status**?
- **Live demo!**
- **Future work**

# EESSI in a nutshell

EESSI



- European Environment for Scientific Software Installations  
(EESSI, pronounced as "easy")
- Collaboration between different partners in HPC community
- Goal:  
**building a common scientific software stack for HPC & beyond**

<https://eessi-hpc.org>

<https://github.com/EESSI>

<https://eessi.github.io/docs/pilot>

 @eessi\_hpc

# EESSI partners & interested parties

Founding partners:



UNIVERSITY OF TWENTE.



**DELL** Technologies

Extensive interest from (HPC) community:



UiO : University of Oslo

# HPC.NRW

# Motivation

EESSI

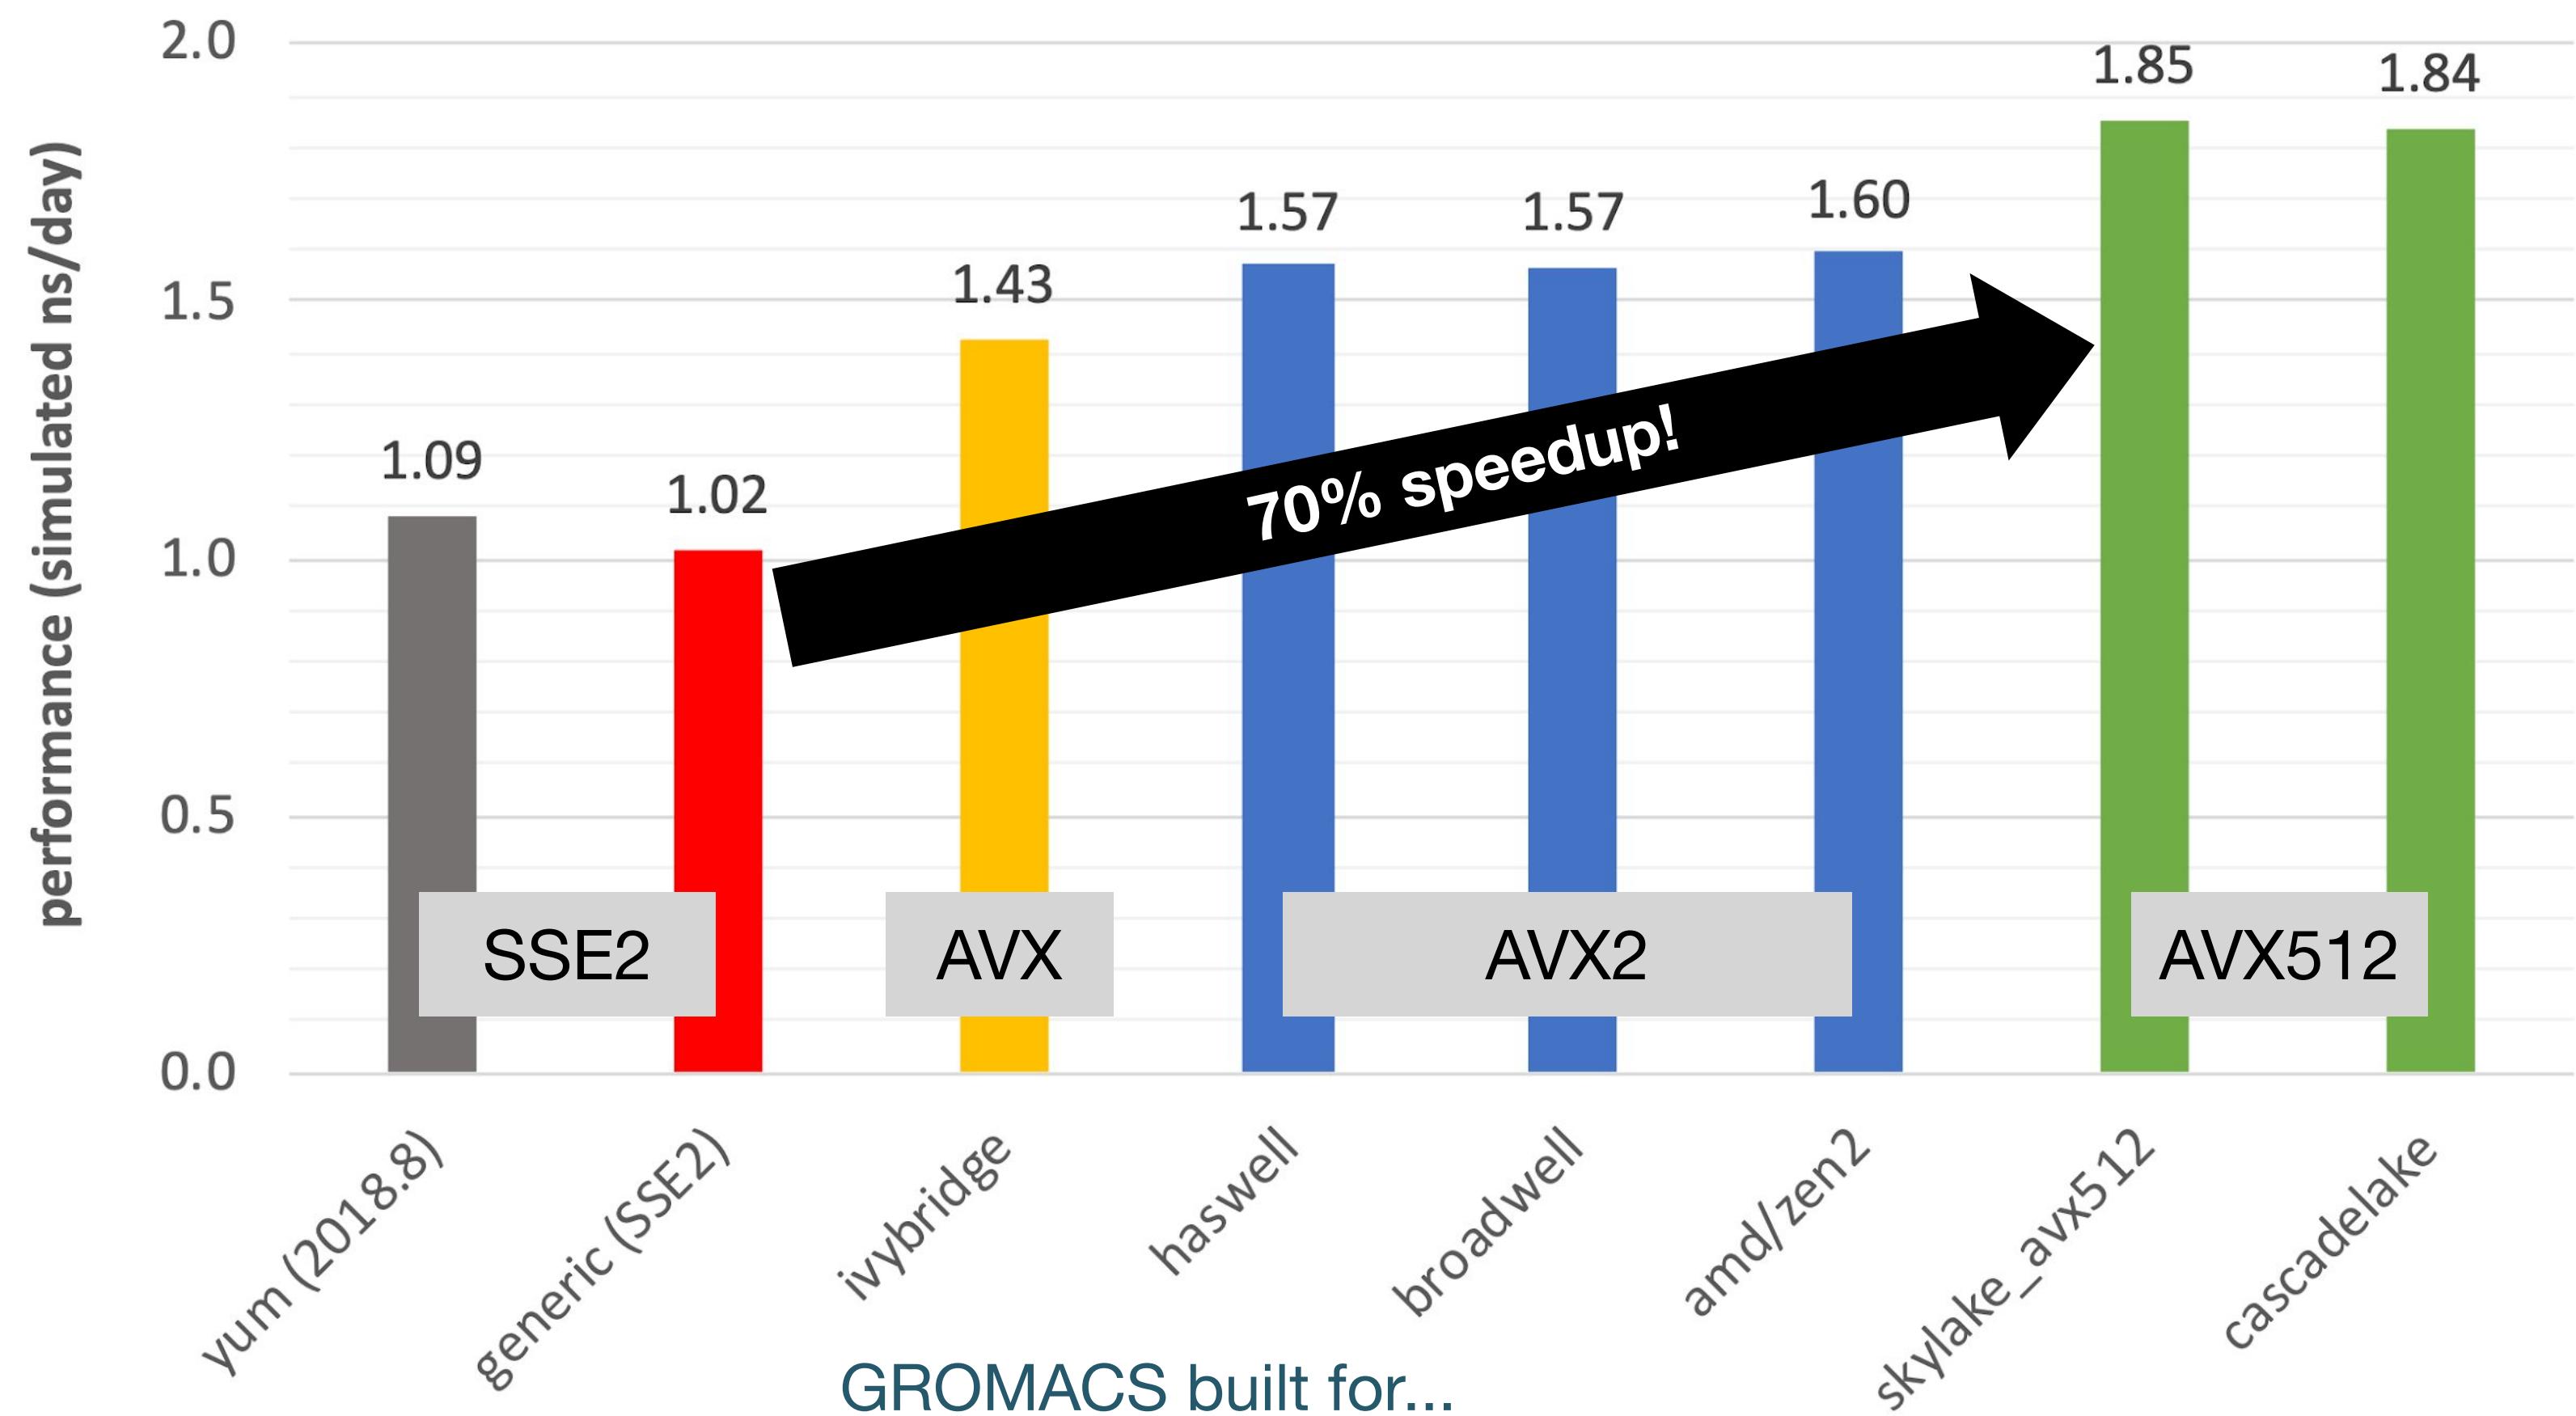


- In short: **more users, more apps, more hardware, more infrastructure**
  - More scientists need to run large computations
  - Explosion of open source scientific software in recent years
  - Increasing variety in CPUs: Intel, AMD, Arm, POWER, RISC-V
  - Various types of accelerators: NVIDIA & AMD GPUs, Intel Xe, ...
  - Rise of the cloud: Amazon EC2, Microsoft Azure, Google, Oracle, ...
- In stark contrast: available manpower in HPC support teams...

# Keeping the P in HPC

- Software should be optimised for the system it will run on
- Impact on performance is often significant for scientific software

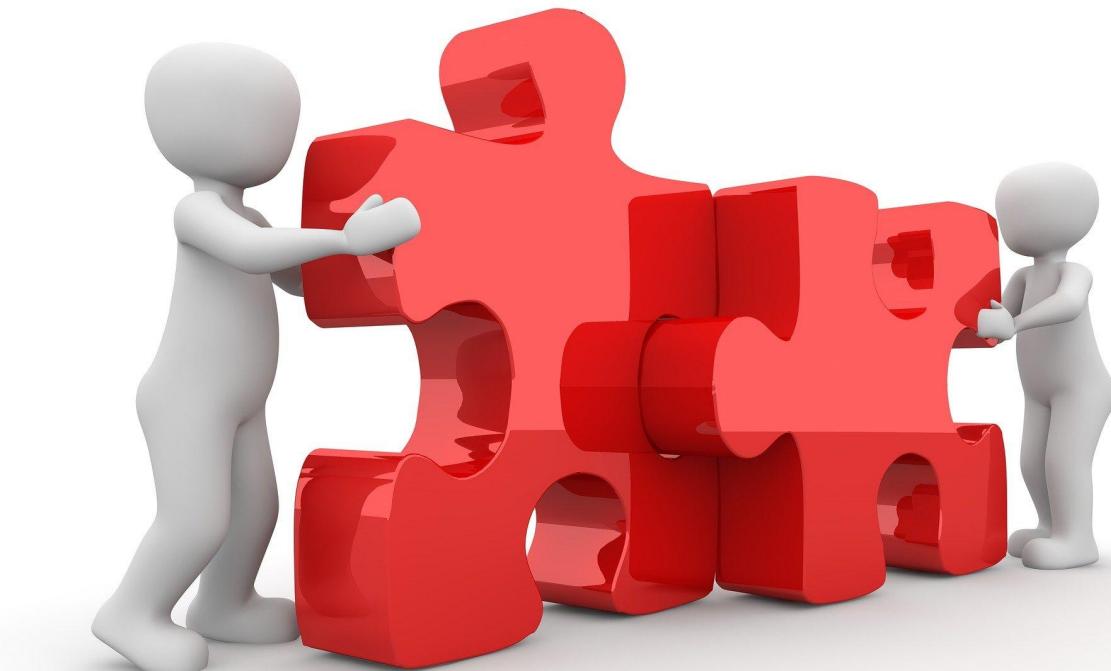
- Example: GROMACS 2020.1  
(PRACE benchmark, Test Case B)
- Metric: (simulated) ns/day,  
higher is better
- Test system: dual-socket  
Intel Xeon Gold 6420  
(Cascade Lake, 2x18 cores)



# Scope & goals



- **Shared repository of scientific software installations**
- **Collaborate**, avoid duplicate work across HPC sites
- Uniform way of providing software to researchers
- Should work on a **variety of systems**:
  - Any Linux distribution, macOS, Windows (via WSL)
  - From laptops and personal workstations to HPC clusters and the cloud
  - Support for different CPUs, interconnects, GPUs, etc
- Focus on **performance, automation, testing, collaboration**



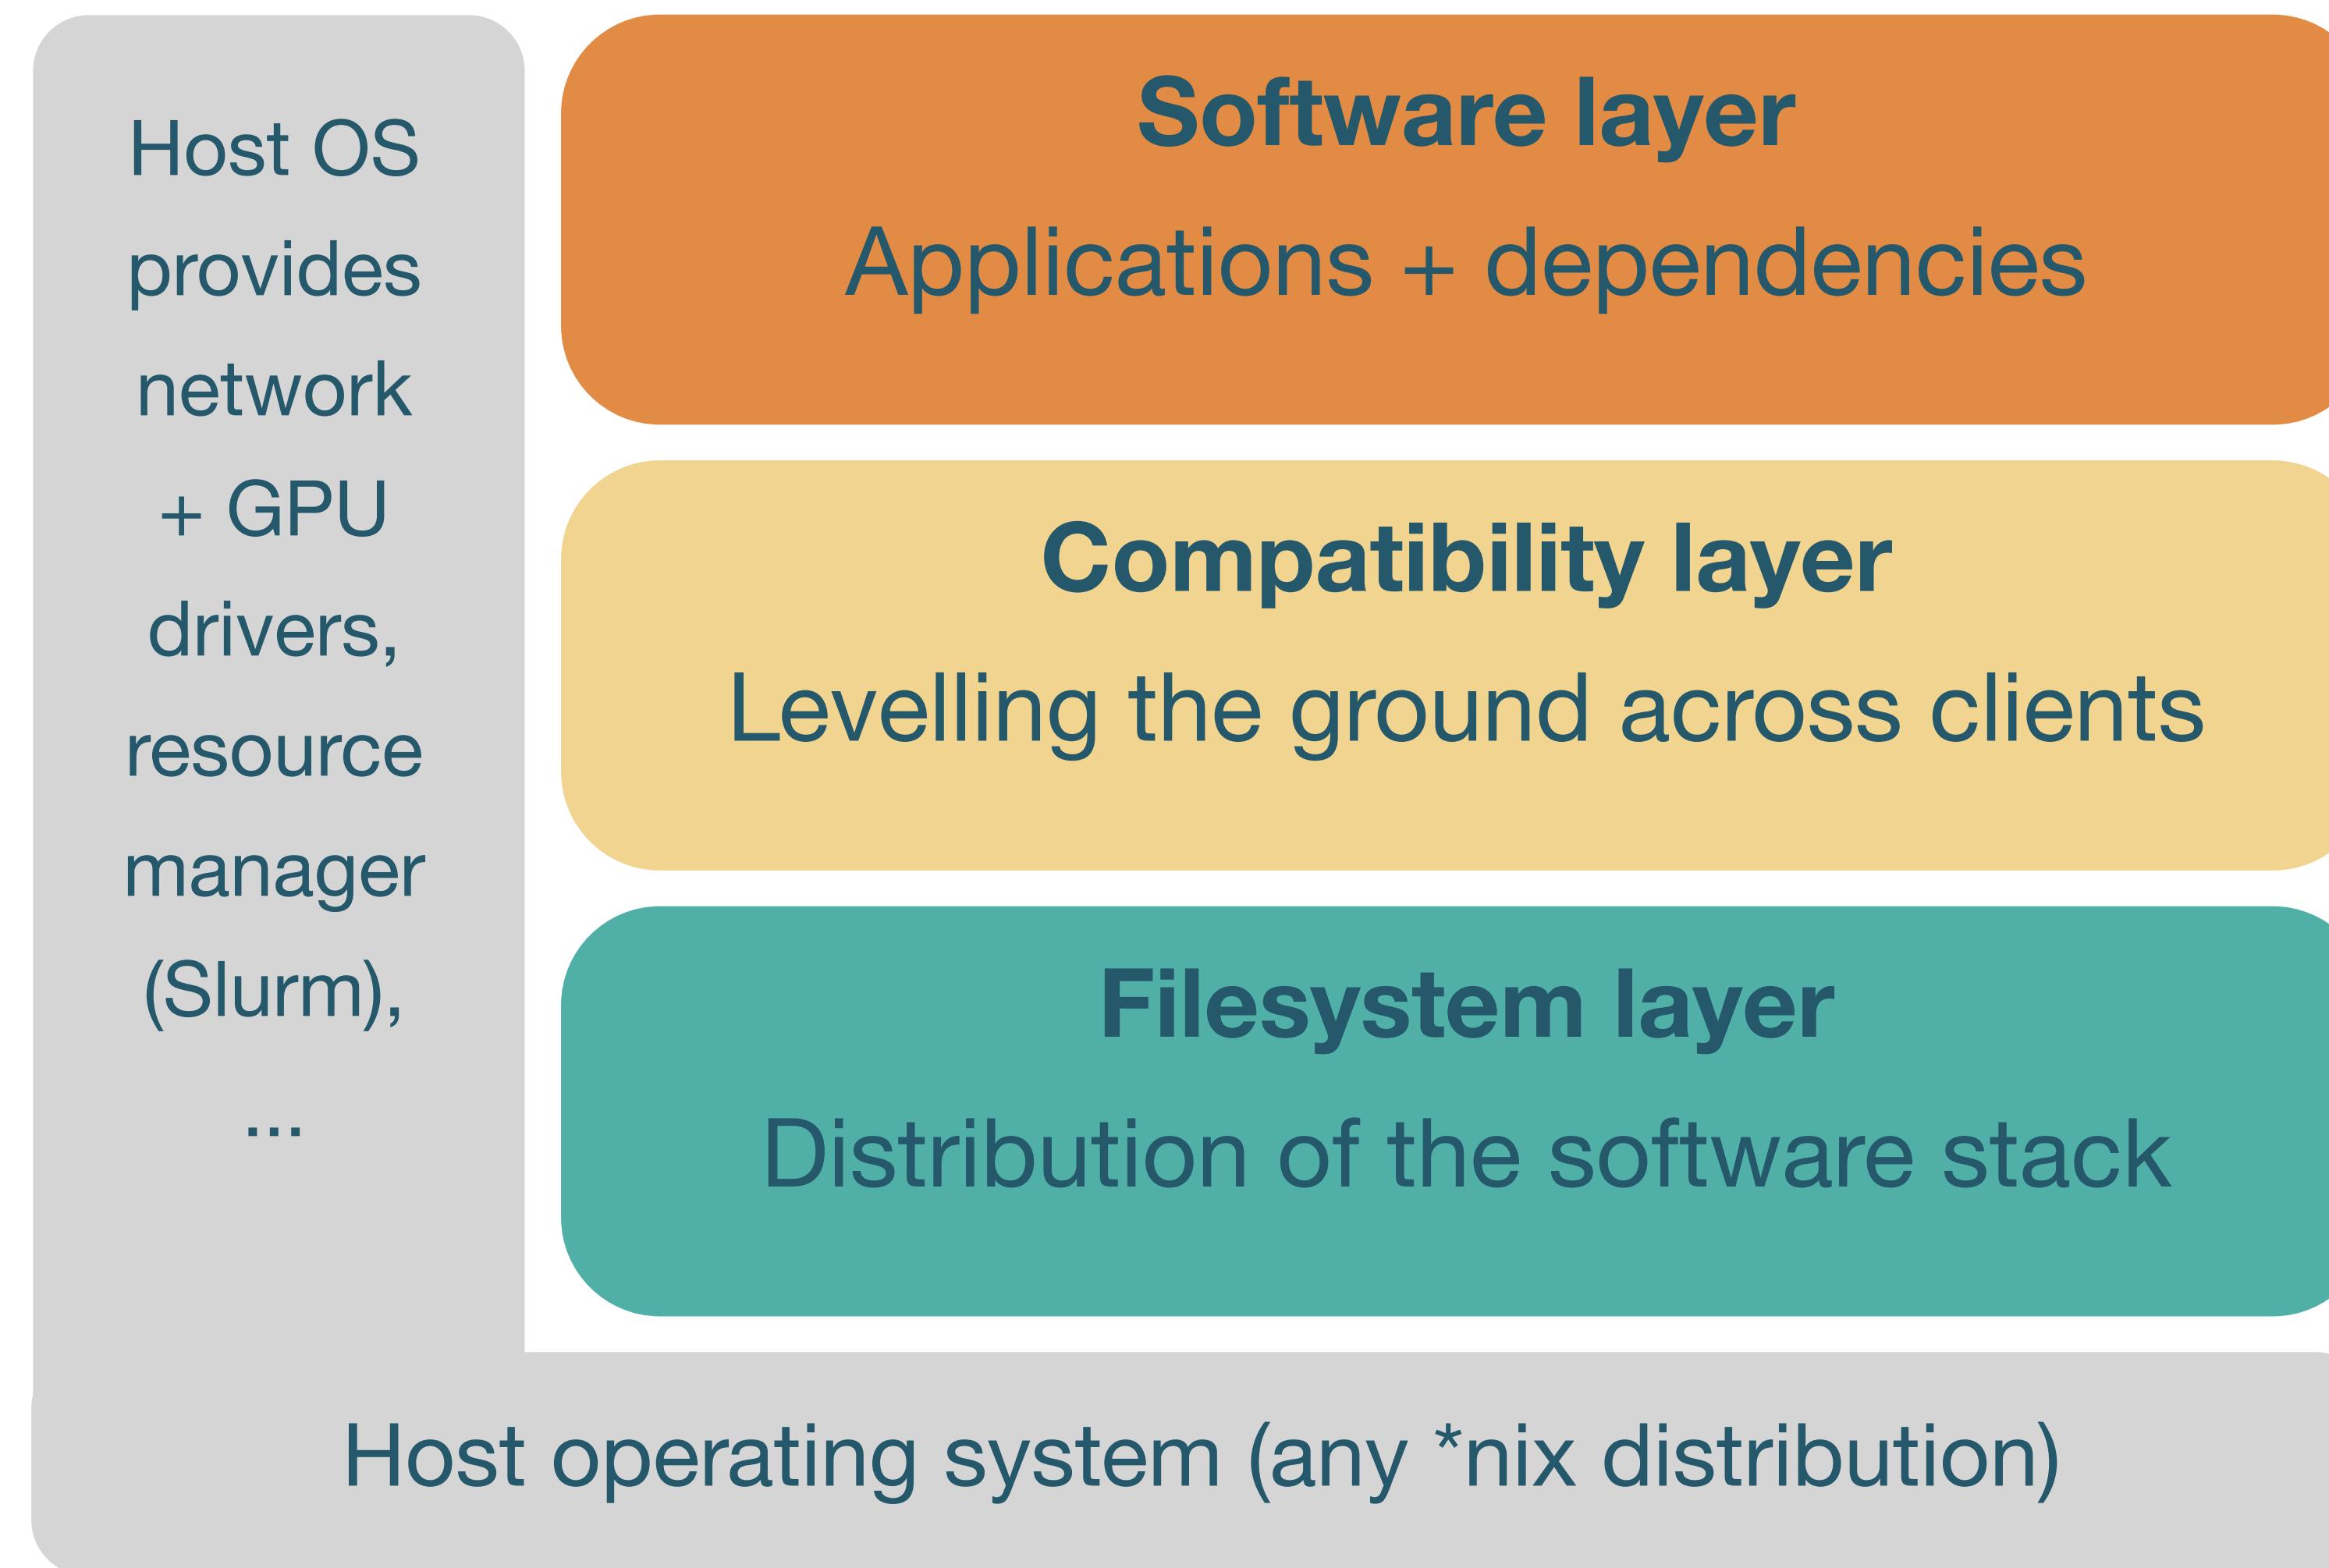
# Inspiration for EESSI



- EESSI concept is **heavily** inspired by Compute Canada software stack
- Shared across 5 major national systems in Canada + a bunch of smaller ones
- **3 layers:** CernVM-FS / ~~Nix~~ Gentoo Prefix / EasyBuild + Lmod
- See paper by Maxime Boissonneault & co at PEARC'19 (PDF available [here](#))

*“Providing a Unified Software Environment for Canada’s National Advanced Computing Centers”*
- See also Maxime’s talk at 5th EasyBuild User Meeting ([slides](#) - [recorded talk](#)) and the Compute Canada [documentation](#)

# High-level overview of EESSI project



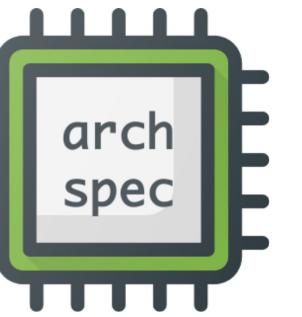
# EESSI is powered by FOSS (1/2)

EESSI



- Installation tool for scientific software
- Optimises for build host (by default)
- Supports over 2,000 different pkgs

<https://easybuild.io/eum>



- Python library
- Detects processor type
- Check compatibility with host CPU

<https://github.com/archspec>



- Environment modules tool (in Lua)
- Intuitive access to software installations
- Multiple versions side-by-side

<https://lmod.readthedocs.io>



gentoo

<https://wiki.gentoo.org/wiki/Project:Prefix>

- Gentoo: Linux distribution, installs from source
- Prefix subproject: **install packages in <prefix>**
- Supports x86\_64, Arm64, POWER, ...
- Supports both Linux and macOS



CernVM-FS

<https://cernvm.cern.ch/fs>

- Software distribution service (software *installations*, not packages!)
- Scalable, read-only, globally distributed filesystem
- Served by web servers (HTTP only), no firewall issues
- Originally build for Large Hadron Collider (LHC) project at CERN



- Regression testing framework for HPC
- Tests are implemented as Python classes
- Verify correctness
- Evaluate performance

<https://reframe-hpc.rtfd.io>

# EESSI is powered by FOSS (2/2)

EESSI



ANSIBLE

- Tool for automation and configuration management
- Using “playbooks” (YAML)
- **Used to automate deployment of filesystem and compatibility layer**

<https://www.ansible.com>



- Singularity: popular container runtime for HPC
- Own container image format
- Also consumes Docker containers
- **Used to fully control build environment for compatibility and software layer + (optionally) to let clients access EESSI**

<https://sylabs.io/singularity>



Terraform

- “Infrastructure as code” tool
- Creating & managing cloud instances
- Declarative configuration files in custom DSL (HashiCorp Configuration Language - HCL)
- **Planning to use this for creating on-demand build/test nodes in AWS/Azure/...**

<https://www.terraform.io>



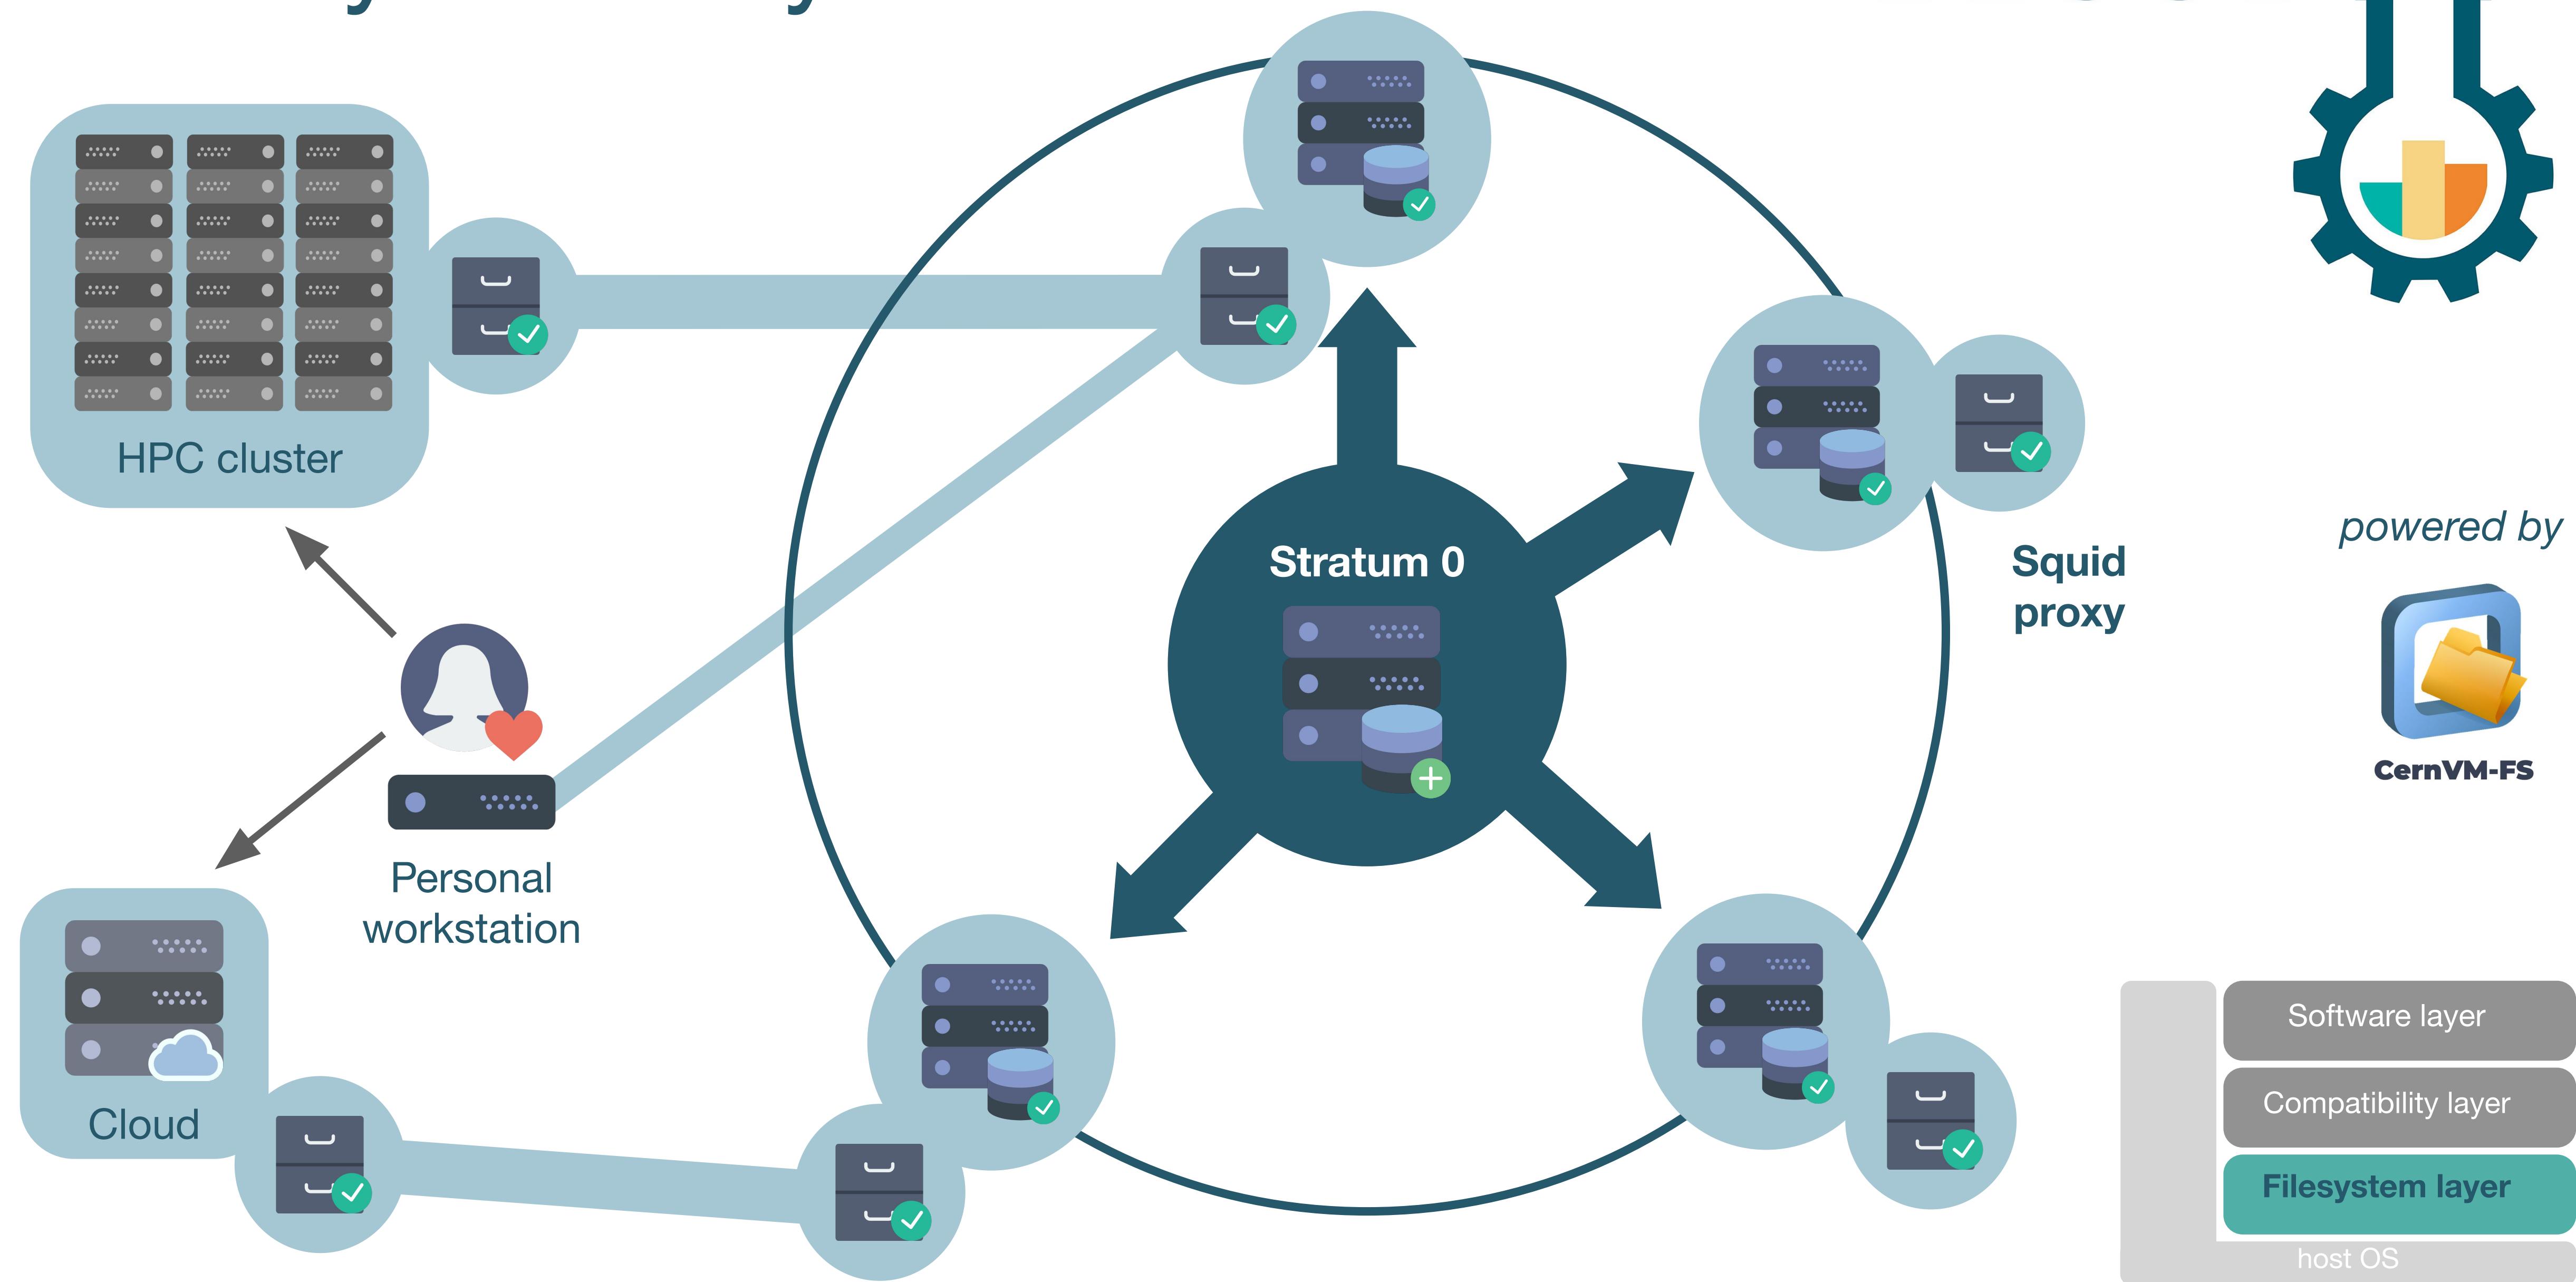
CLUSTER IN  
THE CLOUD

- CLI tool to easily create disposable Slurm clusters in the cloud
- Supports AWS, Oracle, Google cloud (Azure not yet supported)
- Leverages Ansible and Terraform in the background
- **Planning to use this to set up (heterogenous) Slurm clusters for building and testing software**

<https://github.com/clusterinthecloud>

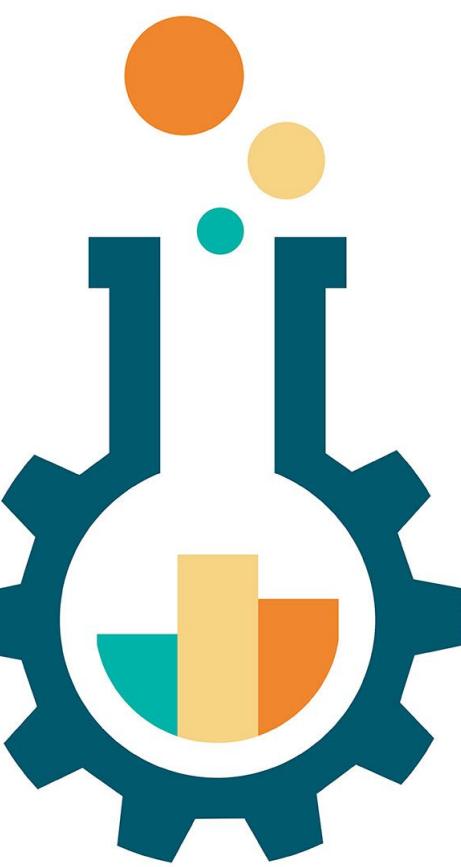
# Filesystem layer

EESSI



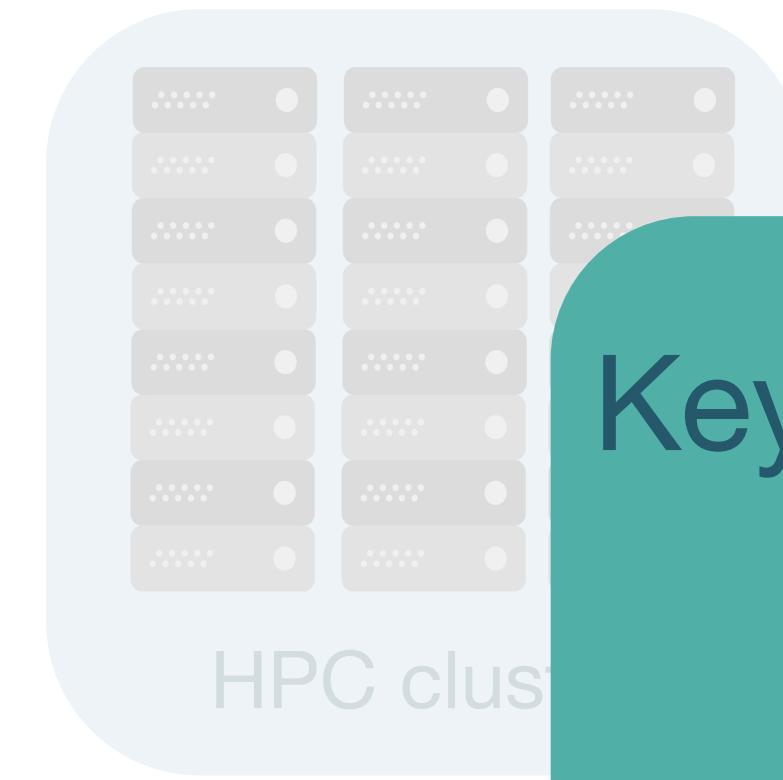
# Filesystem layer

EEESSI

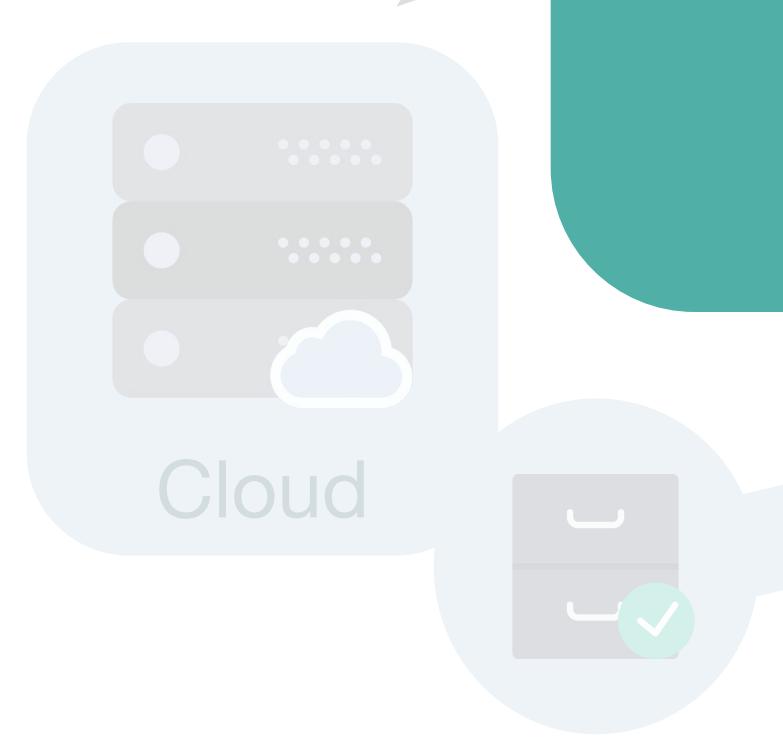


## Key messages:

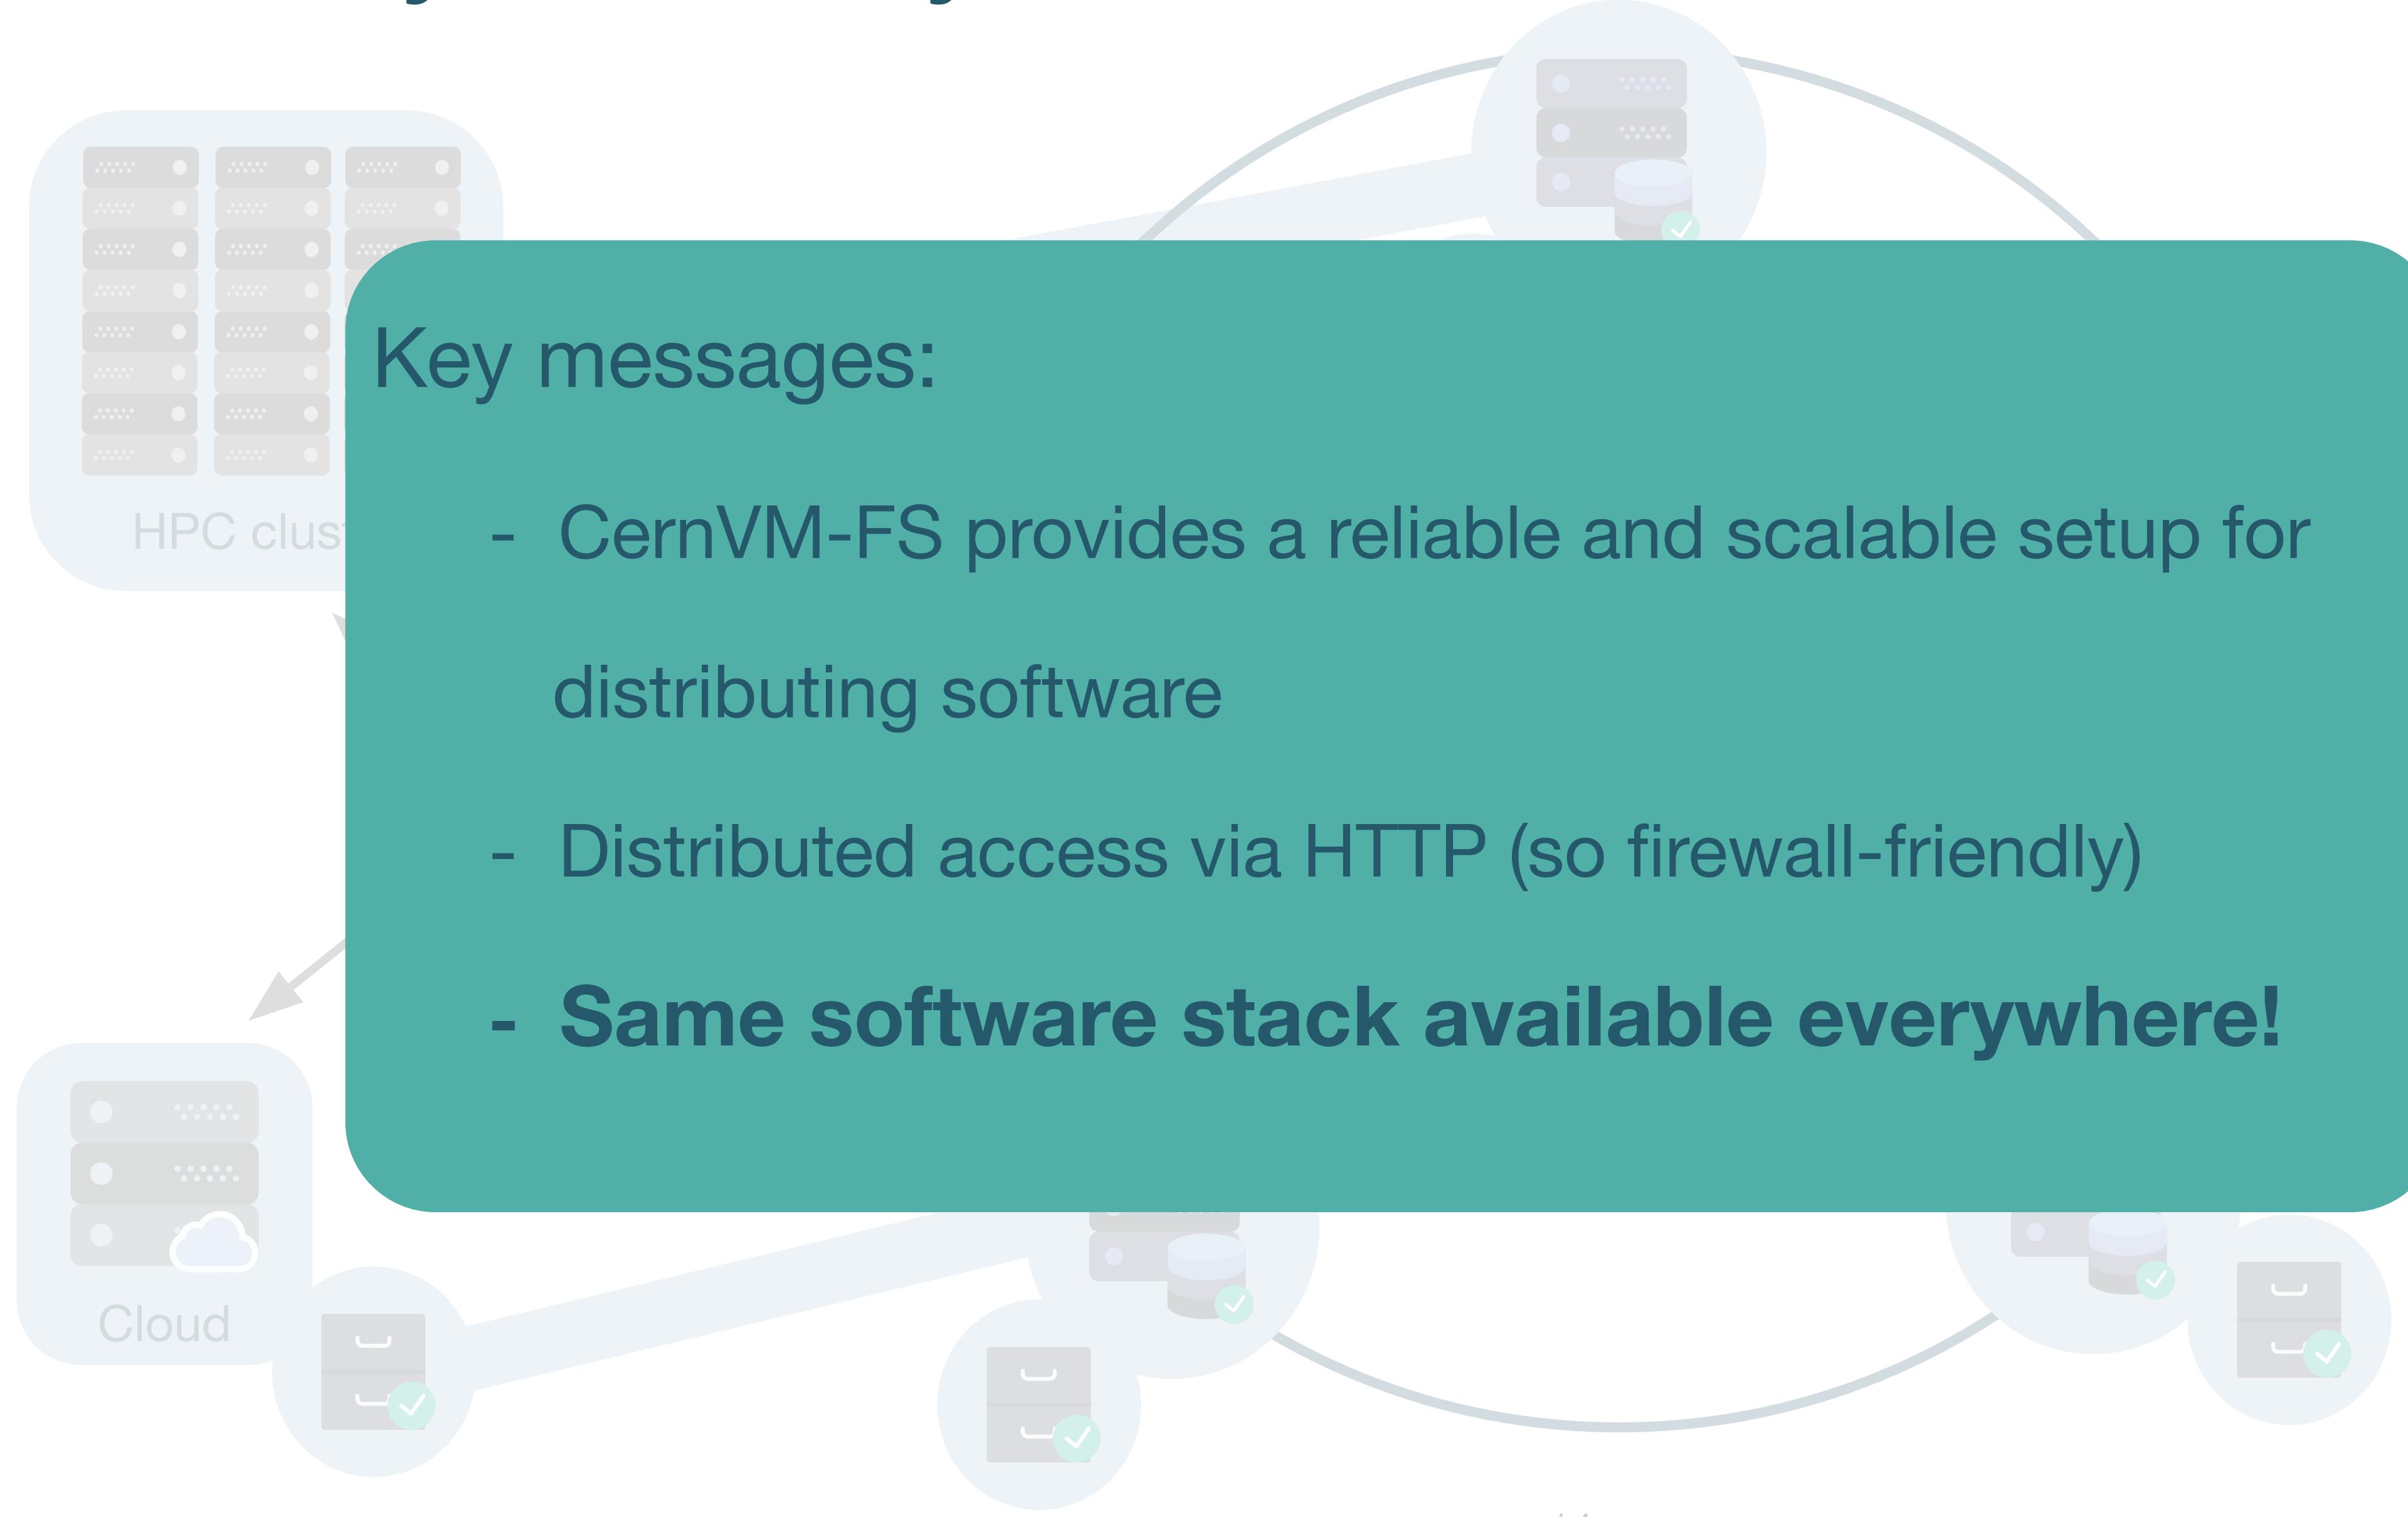
- CernVM-FS provides a reliable and scalable setup for distributing software
- Distributed access via HTTP (so firewall-friendly)
- **Same software stack available everywhere!**



HPC cluster



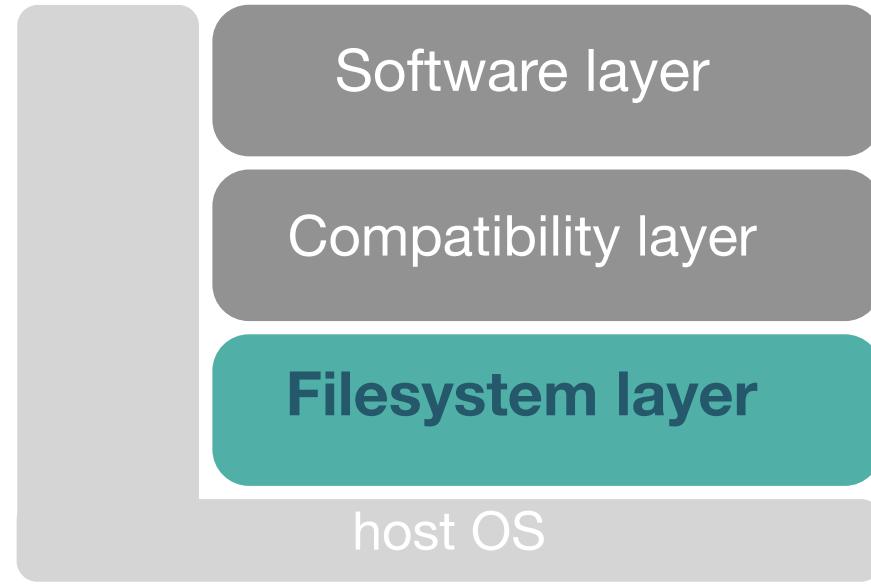
Cloud



powered by



CernVM-FS



Software layer

Compatibility layer

**Filesystem layer**

host OS

# Compatibility layer

- Gentoo Prefix installation
- Set of tools & libraries installed in non-standard location
- Limited to low-level stuff, incl. glibc. No kernel or drivers.
- Only targets a supported processor **family** (x86\_64, Arm64, ppc64le)
- **Levels the ground for different client operating systems** (Linux distros, macOS)
- Currently in pilot repository:

/cvmfs/pilot.eessi-hpc/2020.12/compat/linux/aarch64

/cvmfs/pilot.eessi-hpc/2020.12/compat/linux/ppc64le

/cvmfs/pilot.eessi-hpc/2020.12/compat/linux/x86\_64



EESSI

*powered by*



Software layer

Compatibility layer

Filesystem layer

host OS

# Software layer



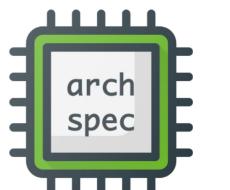
ESSI

- Provides scientific software applications, libraries, and dependencies
- **Optimised for specific CPU microarchitectures** (Intel Haswell, ...)
- **Leverages libraries from compatibility layer** (not from host OS)
- Installed with EasyBuild, incl. environment module files
  - Intention to start using Easystack files soon
- Lmod environment modules tool is used to access installations
- Different subdirectories/trees: one per CPU microarchitecture
- **Best subdirectory for host is picked automatically** via archspec

*powered by*



Lmod



Software layer

Compatibility layer

Filesystem layer

host OS

# Software build and ingestion procedure

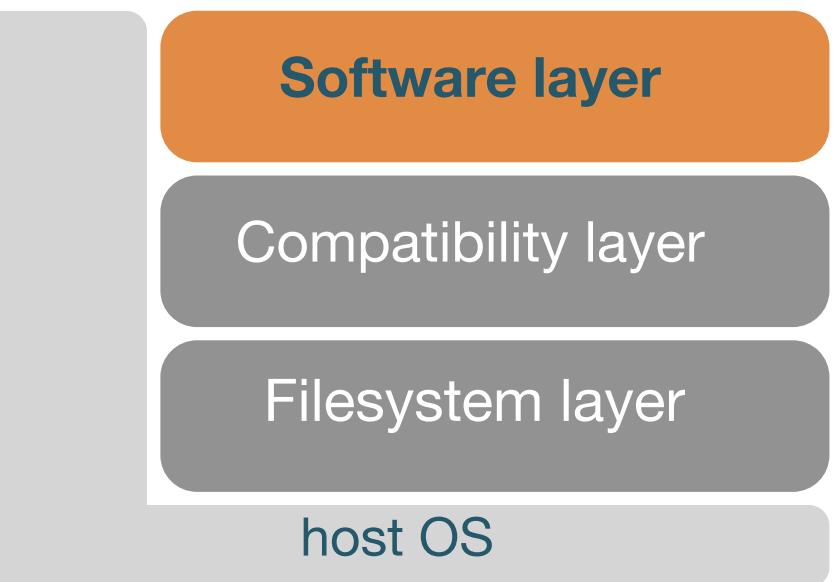


- Build machines with specific microarchitectures
- Singularity build containers
  - **No root privileges** required
  - FUSE mount for CVMFS and **writable overlay** (fuse-overlayfs)
  - **Controlled and minimal** build environment
- Manually start **build script** that runs EasyBuild
- Tarballs of the overlay's upper dir are manually ingested on Stratum 0
- We want to **automate** this whole process using GitHub PRs
  - Any experiences with such approaches?

*powered by*  
 **easybuild**



**fuse-overlayfs**



# Current status: pilot repository

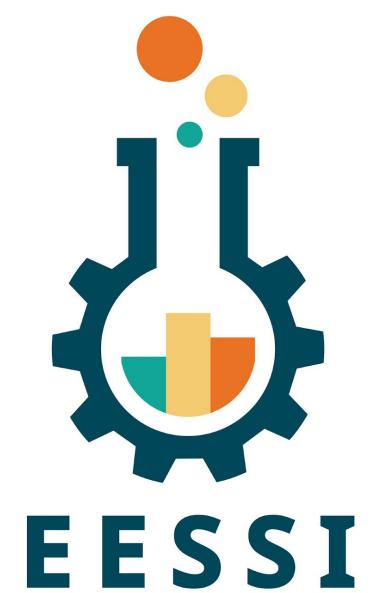


- Ansible playbooks, scripts, docs at <https://github.com/eessi>
- CernVM-FS: Stratum 0 @ Univ. of Groningen + two Stratum 1 servers
- Compatibility layer for x86\_64, ppc64le, aarch64 (only Linux clients, for now)
- Software (CPU-only): Bioconductor, GROMACS, OpenFOAM, TensorFlow
- Hardware targets:
  - x86\_64/generic, intel/haswell, intel/skylake\_avx512, amd/zen2
  - aarch64/generic, aarch64/graviton2, aarch64/thunderx2

**NOT FOR  
PRODUCTION USE!**

Try it yourself: <https://eessi.github.io/docs/pilot>

# From zero to science in three steps



1. Access the EESSI CernVM-FS repo

- Native installation of the CernVM-FS client  
**(requires admin privileges)**



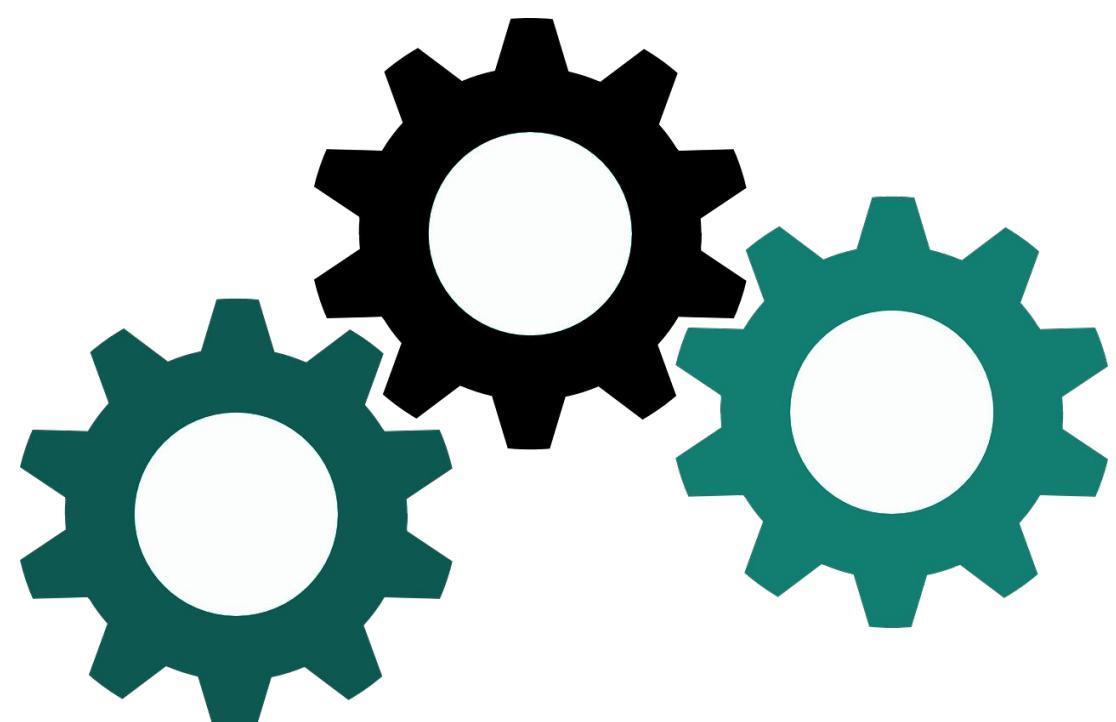
2. Source the EESSI init script

Detect your microarchitecture,  
find the right software tree,  
set up your environment.



3. Compute!

Load the module(s) that you need, and start running!



# Step 1: Access the EESSI repository

## Option 1 (example):

native CernVM-FS installation on fresh (x86\_64) RHEL 8.2 system



```
# install CernVM-FS client (see https://cernvm.cern.ch/fs/)
sudo yum install -y https://ecsft.cern.ch/dist/cvmfs/cvmfs-release/cvmfs-release-latest.noarch.rpm
sudo yum install -y cvmfs

# install CernVM-FS configuration files for EESSI repositories (see https://github.com/EESSI/filesystem-layer)
wget https://github.com/EESSI/filesystem-layer/releases/download/v0.2.3/cvmfs-config-eessi-0.2.3-1.noarch.rpm
sudo yum install -y cvmfs-config-eessi-0.2.3-1.noarch.rpm

# create local CernVM-FS configuration file (direct access, no proxy; 10GB for CernVM-FS cache)
sudo bash -c "echo 'CVMFS_HTTP_PROXY=DIRECT' > /etc/cvmfs/default.local"
sudo bash -c "echo 'CVMFS_QUOTA_LIMIT=10000' >> /etc/cvmfs/default.local"

# set up CernVM-FS
sudo cvmfs_config setup

# access EESSI pilot repository
ls /cvmfs/pilot.eessi-hpc.org/2020.12
```

# Step 1: Access the EESSI repository

**Option 2** (example, see <https://eessi.github.io/docs/pilot>):

use Singularity to run Docker container to access EESSI

```
# configure Singularity (bind mounts + home directory)
mkdir -p /tmp/$USER/{var-lib-cvmfs,var-run-cvmfs,home}
export SINGULARITY_BIND="/tmp/$USER/var-run-cvmfs:/var/run/cvmfs,/tmp/$USER/var-lib-cvmfs:/var/lib/cvmfs"
export SINGULARITY_HOME="/tmp/$USER/home:/home/$USER"

# values to pass to --fusemount (EESSI config + pilot repositories)
export EESSI_CONFIG="container:cvmfs2 cvmfs-config.eessi-hpc.org /cvmfs/cvmfs-config.eessi-hpc.org"
export EESSI_PILOT="container:cvmfs2 pilot.eessi-hpc.org /cvmfs/pilot.eessi-hpc.org"

# minimal Docker container from Docker Hub (includes CernVM-FS + EESSI configuration files)
export DOCKER_IMAGE="docker://eessi/client-pilot:centos7-$(uname -m)"

# start shell in Singularity container (ignore the scary looking 'setxattr' warnings, they're harmless)
singularity shell --fusemount "$EESSI_CONFIG" --fusemount "$EESSI_PILOT" $DOCKER_IMAGE

# access EESSI pilot repository
ls /cvmfs/pilot.eessi-hpc.org/2020.12
```



# Step 2: source the EESSI init script



```
# source the EESSI init script to set up your environment
$ source /cvmfs/pilot.eessi-hpc.org/2020.12/init/bash
Found EESSI pilot repo @ /cvmfs/pilot.eessi-hpc.org/2020.12!
Found Lmod configuration file at /cvmfs/pilot.eessi-hpc.org/2020.12/software/x86_64/intel/haswell/.lmod/lmodrc.lua
Initializing Lmod...
Prepending /cvmfs/pilot.eessi-hpc.org/2020.12/software/x86_64/intel/haswell/modules/all to $MODULEPATH...
Environment set up to use EESSI pilot software stack, have fun!

[EESSI pilot 2020.12] $ echo $EESSI_PREFIX
/cvmfs/pilot.eessi-hpc.org/2020.12

[EESSI pilot 2020.12] $ echo $EESSI_SOFTWARE_SUBDIR
x86_64/intel/haswell
```

# Step 3: load your modules, and go!



```
# check which modules are available
[EESSI pilot 2020.12] $ module avail gromacs

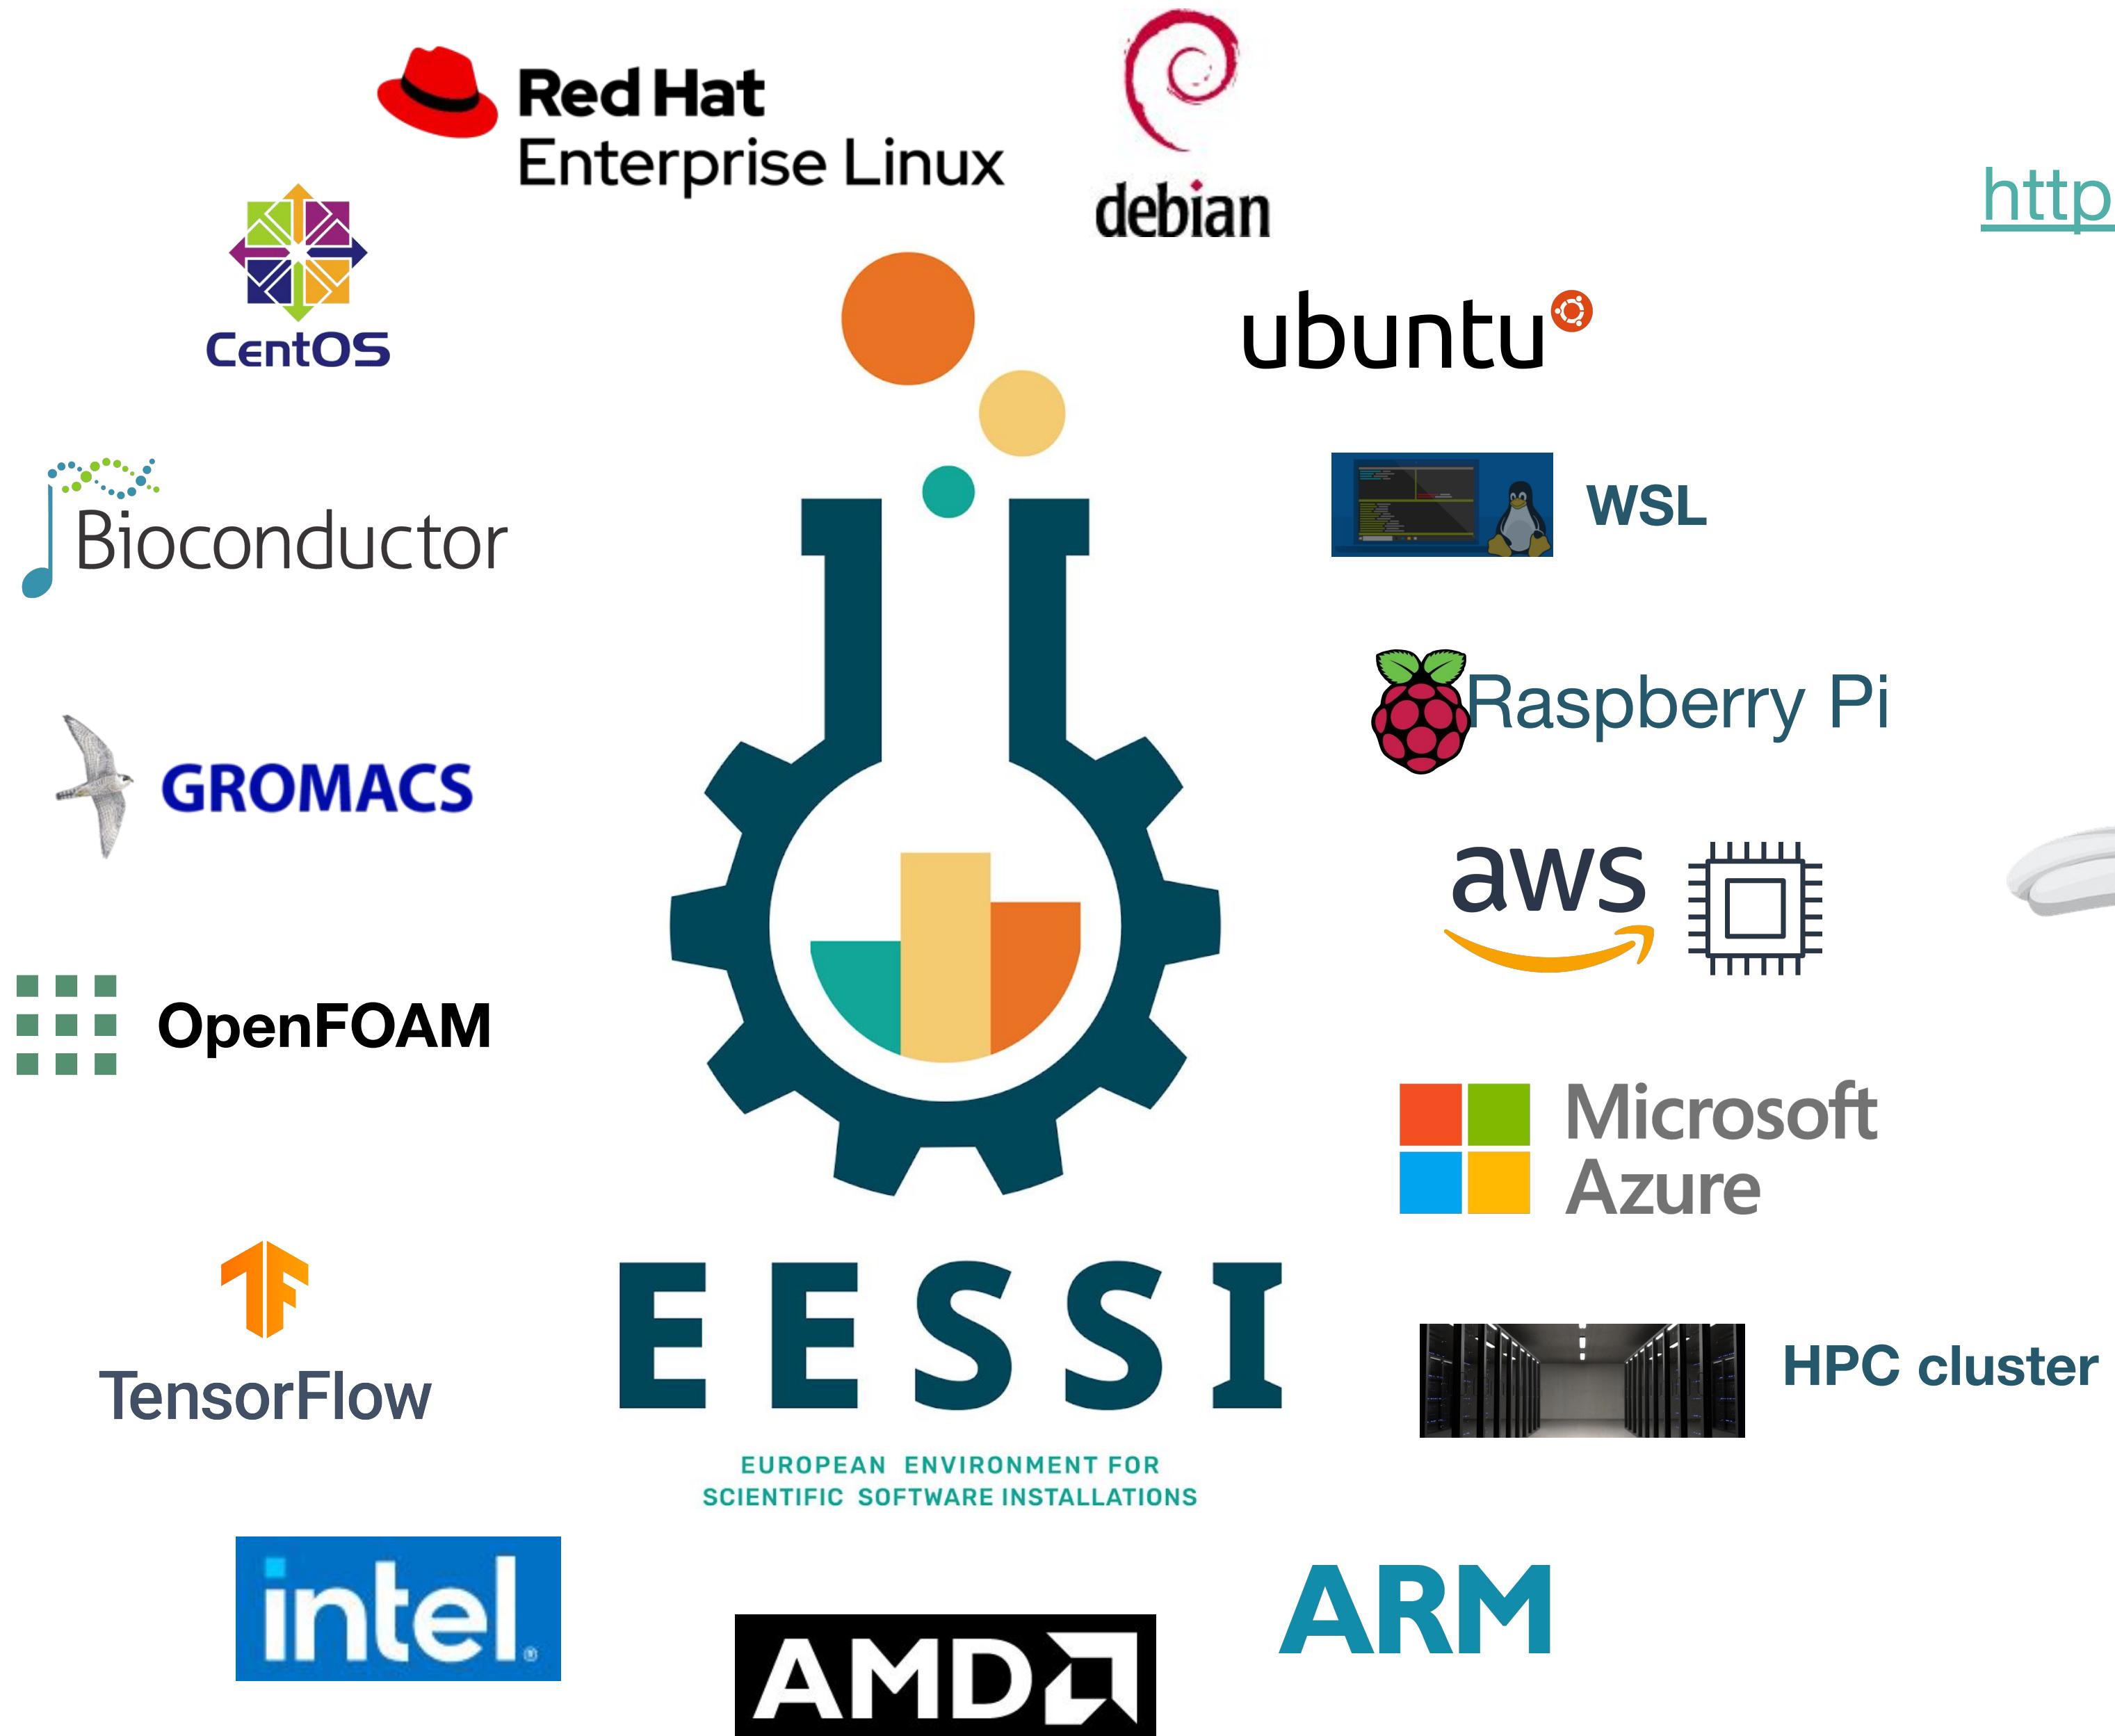
----- /cvmfs/pilot.eessi-hpc.org/2020.12/software/x86_64/intel/haswell/modules/all -----
GROMACS/2020.1-foss-2020a-Python-3.8.2

# load the module(s) for the software you want to use
[EESSI pilot 2020.12] $ module load GROMACS

# ready to compute!
[EESSI pilot 2020.12] $ gmx mdrun -s ion_channel.tpr -maxh 0.50 -resethway -noconfout -nsteps 1000
```

# Demo time!

<https://github.com/ESSI/eessi-demo>



# Demo time!

<https://github.com/ESSI/eessi-demo>

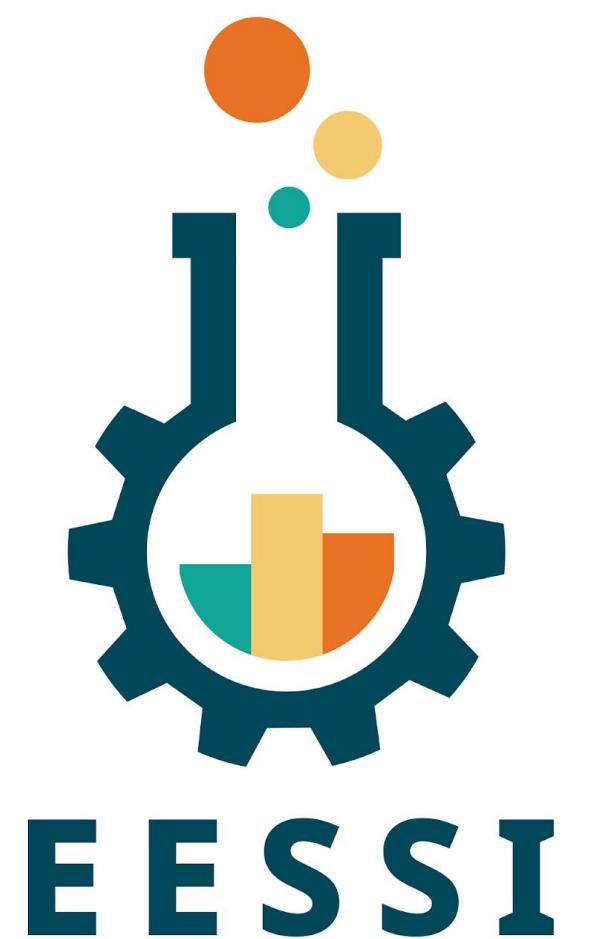


Machine type	Operating system	Architecture	CPU
Laptop	Windows / Ubuntu	x86_64	Intel Core i7-8665U (Whiskey Lake)
VM @ AWS	CentOS 7	aarch64	Graviton2
Openstack VM @ OSU Open Source Lab	CentOS 7	ppc64le	IBM POWER9
Raspberry Pi	Raspbian 10 (Singularity: CentOS 7)	aarch64	Cortex-A72



# CernVM-FS for EESSI / HPC systems

- CernVM-FS is a perfect solution for our filesystem layer
- Lots of HPC sites are already interested
  - Currently, most sites are (still) using/playing with the container
- Ideally, they will be installing the client on their nodes
- Main concern for some sites: **not having a full copy of the software stack**
  - Some software may be unavailable during network outage
  - Recommendations? Using private/internal Stratum 1 servers?



*powered by*



# Future work

- Further improve pilot EESSI repository (monthly revisions)
- Identify problems, and fix them...
- More **automation** (Ansible, Terraform, ...) and **testing** (ReFrame + GitHub Actions)
- Also support macOS / ~~POWER~~ / GPUs, add more software
- Let developers of scientific software validate the installation of *their* software
- Solicit more manpower, get project funded to make it sustainable
- Set up a consortium, and change the “European” in our name
- Work towards **production** setup...





Website: <https://www.eessi-hpc.org>

**Join our mailing list & Slack channel**  
<https://www.eessi-hpc.org/join>

Documentation: <https://eessi.github.io/docs>

GitHub: <https://github.com/eessi>

Twitter: [@eessi\\_hpc](https://twitter.com/eessi_hpc)

Monthly online meetings (first Thursday, 2pm CET)

CernVM-FS tutorial @ 6th EasyBuild User Meeting (Jan 25-29, 2021):  
<https://cvmfs-contrib.github.io/cvmfs-tutorial-2021/>